

The Benjamin Button Problem

Ashby's Constraint and the Agentic Enterprise

Rajesh Iyer

iyer70@gmail.com

Abstract

The industry tried to be born old—deploying agents before earning the governance to sustain them. For eighteen months, the narrative held that LLMs reason like humans, therefore agents everywhere. That story skipped the hard parts: governance, credibility, institutional risk, coordination cost. Enterprises didn't reject agents. They priced them correctly. This paper argues for aging backwards: beginning with GenAI-augmented automation, graduating to agentic patterns, and instantiating agents only as a reward for earned governance. The theoretical foundation combines Ashby's Law of Requisite Variety with Blackman's insight on stability: you cannot govern a high-variety system with a low-variety regulator, and low-frequency governance cannot stably control high-frequency agent behavior. Together, these constraints define a regime boundary—below which deterministic governance suffices, above which agentic governance becomes mandatory. Crucially, GenAI itself enables the crossing: by augmenting governance at computational frequency, it solves the frequentist problem that makes agentic governance possible.

1. The Goal Is Acceleration

Enterprises want the value GenAI makes possible—faster decisions, semantic flexibility, judgment at scale. The question is not whether to pursue agentic capability, but how to sequence its adoption. This paper argues that the fastest sustainable path is counterintuitive: start with automation, earn governance capability, and treat agent instantiation as a reward rather than a starting point. The industry's mistake was not ambition. It was sequence.

2. The Mistake the Industry Made

For eighteen months the industry sold a story: "LLMs reason like humans, therefore agents everywhere."

That story skipped the hard parts:

- governance
- credibility
- institutional risk
- coordination cost

Enterprises didn't reject agents. They priced them correctly. When they did, the math flipped.

The fundamental error was a conflation of three distinct capabilities:

Capability	Source	Gov. Cost
Reasoning	LLMs provide this	Low
Autonomy	Automation handles this	Known
Persistence	Where risk explodes	Superlinear

LLMs gave us reasoning. Automation already handled bounded autonomy. Persistence—stateful, delegating,

identity-bearing agents—is where governance cost explodes.

Enterprises said: "*Give us reasoning inside systems we already know how to govern.*"

That's not cowardice. That's maturity.

3. The Hidden Truth

Agentic capability is real. Agent instantiation is expensive.

Most enterprise work does not require:

- persistent identity
- autonomous delegation
- long-lived state
- recursive planning

It requires:

- judgment compression
- language-to-action translation
- ambiguity resolution
- decision support under constraint

That's agentic *behavior*, not agents. And agentic behavior fits beautifully inside automation.

4. Definitions

Two terms must be distinguished, as their conflation drives most enterprise AI failures:

Agentic (adjective): Exhibiting autonomy, goal-pursuit, and action-taking capacity. A property that systems may possess to varying degrees.

Agent (noun): A discrete, stateful, persistent entity with identity, goals, and delegation capacity. A reification of

agentic properties into a bounded, accountable unit.
Agentic is a property. Agents are a reification—expensive, powerful, dangerous.

A system may be agentic without deploying agents.

5. Why Automation Wins Now

Automation with GenAI is more transformative than agents right now.

Why? Because it:

- turns brittle workflows into semantic ones
- upgrades rules into judgment
- makes every system “soft” without making it autonomous
- removes humans from toil without removing humans from intent

That's not step one. That's step zero done correctly.

From a governance perspective, automation:

- inherits existing controls
- produces auditable artifacts
- collapses time-to-coherence
- doesn't demand new trust primitives

Agents, by contrast:

- require new governance layers
- create new failure modes
- surface before credibility is earned

Automation is the easier path because it's the only path that composes with reality today.

6. Ashby and Blackman: The Two Constraints

Two constraints from cybernetics and signal processing govern what is architecturally possible.

6.1 Ashby's Constraint (Variety)

In 1956, British cyberneticist W. Ross Ashby articulated the Law of Requisite Variety in *An Introduction to Cybernetics*:

“Only variety can absorb variety.”

A regulator must possess at least as much complexity as the system it governs. Simple controllers fail against complex systems. The math is unforgiving.

Cybernetics—from the Greek *kybernetes*, “steersman”—is the study of how systems regulate themselves through feedback. Ashby's insight was foundational to control theory, systems biology, and organizational design.

Applied to agents: Governance must match the variety of agent behavior. Deterministic rules cannot enumerate the state space of persistent, delegating agents.

6.2 Blackman's Constraint (Stability)

Ashby tells us governance must have sufficient *variety*. But variety alone is not enough.

In signal processing, R.B. Blackman and J.W. Tukey demonstrated that sharp filters—binary, hard-edged responses—cause ringing and instability when processing high-frequency signals. The Blackman window, introduced in their 1958 work *The Measurement of Power Spectra*, showed that *damped, tapered responses* maintain stability where sharp cutoffs fail.

“Low-frequency governance cannot stably control high-frequency behavior.”

Applied to governance: sharp approve/deny gates, rigid thresholds, and post-hoc review cause instability when the input is high-frequency and high-variance. The response must be damped, adaptive, and frequency-matched.

Humans are low-frequency samplers:

Governance Type	Frequency
Static rules	Once (design time)
Human review	Hours to days
Committee approval	Days to weeks
Audit	Quarterly

Agents operate at computational frequency—milliseconds to seconds. You cannot sample a gigahertz signal with a kilohertz sampler. Sharp, binary governance causes instability under high-frequency input.

Applied to agents: Human-in-the-loop governance fails not because humans are *slow*, but because they are *low-frequency*. Adding more humans doesn't change the frequency band—it adds latency and coordination cost.

6.3 The Combined Constraint

Together, Ashby and Blackman force a conclusion:

- Ashby rules out simple governance
- Blackman rules out human-frequency governance

For persistent agents, only **agentic governance**—governance that operates at computational frequency with sufficient variety—is architecturally stable.

This is not a recommendation. It is a derivation.

7. The Regime Boundary

The two constraints define a **regime boundary**. Informally (we use this notation loosely, as intuition rather than proof):

System	Variety	Frequency
Bounded patterns (E0–E2)	$O(n)$	Synchronous
Persistent agents (E3)	$O(2^n)$	Computational
Deterministic governance	$O(n)$	Static
Human review	$O(\log n)$	Low
Agentic governance	$O(2^n)$	Computational

The notation is illustrative. The point: bounded patterns can be matched by deterministic governance; persistent agents cannot. The gap is exponential in character.

Below the boundary (E0–E2): Variety bounded, frequency matched. Deterministic governance suffices.

Above the boundary (E3): Variety explodes, frequency mismatches. Only agentic governance can close the gap. Crossing the boundary changes what governance is required.

8. The Execution Sequence

E0: GenAI-Augmented Automation

Judgment inside guardrails. Semantic flexibility without autonomy. 90% of enterprise value lives here.

E1: Agentic Patterns

Stateless, tool-using, constrained. No persistent identity, no delegation, no long-lived state.

E2: Agentic Workflows

Conditional, multi-step, orchestrated. State within workflow bounds, not across sessions.

E3: Agent Instantiation

Stateful, persistent, delegating. Rare. High-value. The reward for earned governance.

Anyone skipping to E3 is lighting credibility on fire.

9. The Governance Sequence

If agents must be earned, so must agentic governance:

Level	Characteristics
G0: Static rules	Hard-coded, enumerated
G1: Parameterized	Configurable thresholds
G2: GenAI-assisted	Adaptive, computational frequency
G3: Agentic	Agents governing agents

G2 is the critical transition—it closes the frequency gap while preserving human intent at the policy layer.

10. GenAI as the Unlock

GenAI solves the frequentist problem.

Gov. Type	Frequency	Variety
Static rules	Once	Low
Human review	$\sim 10^{-5}$ Hz	Low
GenAI-assisted	$\sim 10^0$ – 10^2 Hz	High
Agentic governance	$\sim 10^2$ – 10^3 Hz	High

GenAI doesn't just speed up governance. It **changes the frequency band**. That's not incremental. That's a phase transition.

The path across the regime boundary:

1. **GenAI augments execution** (E0–E1)—deterministic governance holds
2. **GenAI augments governance** (G2)—frequency match achieved
3. **Agentic governance emerges** (G3)—both at compute frequency

The second move enables the third. Without GenAI-assisted governance, you cannot cross the regime boundary safely.

11. The Axiom

From the two constraints, the axiom follows:

"Agentic capability that scales requires agentic governance that scales. Humans govern intent."

The axiom does not claim humans scale execution—only intent. Humans set policy. GenAI enforces it at frequency. Agentic governance emerges when both execution and governance mature together.

12. The Six Laws

Law I: Cognitive Decoupling

Agentic systems scale content faster than humans scale credibility. Credibility becomes the bottleneck, not cognition.

Law II: Agentic Overhead

The governance cost of agents grows superlinearly with their capability. Agentic patterns should be preferred; agent instantiation is a last resort.

Law III: Discrete Scaling Walls

Human systems degrade smoothly at the individual level but fail abruptly at coordination thresholds. Agentic capability creates value only at these discontinuities.

Law IV: Feasibility Dominance

Opportunity does not compound without feasibility; it collapses. Governance beats vision in institutional settings.

Law V: Talent Compression

Agentic tools steepen early productivity curves and flatten late-career decay. Hierarchies compress but do not invert. Judgment remains scarce.

Law VI: Success Becomes Invisible

Successful agentic systems shift visibility from behavior to assurance. Fewer dashboards, more attestations. If the AI is still exciting, it's not working yet.

13. The Three Scaling Truths

Truth	Statement
I	Cognition scales with compute
II	Credibility scales with time
III	Governance must scale with capability

Credibility constrains governance *legitimacy*, not enforcement capacity.

14. The Four Corollaries

I: The Governance Gap. The difference between agentic capability growth and human governance capacity is the enterprise AI problem. Measured in risk, not inefficiency.

II: The Credibility Constant. No compute purchases credibility. No headcount parallelizes trust. Only time and judgment compound.


III: The Visibility Shift. Maturity means visibility shifts from behavior to assurance. Dashboards yield to attestations.

IV: The Restraint Advantage. In agent inflation, advantage belongs to those who know when agentic patterns suffice.

15. The Benjamin Button Regression

The industry's error: attempting to be born old.

The correction: age backwards. Both execution *and* governance must mature together. Each governance improvement unlocks a step-change in execution capability.



Each governance step enables a discrete jump in execution capability. The critical insight: **crossing the regime boundary requires G2**. The step-change from E2 to E3 is only possible once governance operates at computational frequency. Without that unlock, agent instantiation is architecturally unstable.

16. Conclusion

We are not "back to automation." **Agents are the reward for earned governance**—and GenAI makes that earning possible by augmenting governance to computational frequency.

These are institutional scaling laws extending Ashby and Blackman to AI governance. They explain why agentic AI is powerful and dangerous, why most initiatives should die, and why governance is leverage, not overhead.

In one sentence: Agentic capability does. Agentic governance governs. Humans govern intent.

In five words: Capability. Governance. Intent. Credibility. Time.

References

- [1] Ashby, W.R. (1956). *An Introduction to Cybernetics*. Chapman & Hall, London.
- [2] Blackman, R.B. and Tukey, J.W. (1958). *The Measurement of Power Spectra*. Dover Publications, New York.
- [3] Wiener, N. (1948). *Cybernetics: Or Control and Communication in the Animal and the Machine*. MIT Press, Cambridge.

- [4] Board of Governors of the Federal Reserve System (2011).
SR 11-7: Guidance on Model Risk Management. Federal Reserve, Washington.

corpXiv preprint. December 2025.